

**Author: Kunjan Mhaske**

---

## **FASHION-MNIST MODEL**

---

### **Dataset: Fashion-MNIST**

Fashion-MNIST is a dataset of Zalando's article images which consists

- Training set of 60,000 examples
- Test set of 10,000 examples
- Each example is a 28 x 28 grayscale image: 784 pixels valued between 0 and 255
- Each associated with a label from 10 classes
- Each row is a separate image
- Column 1 is the class label.
- Remaining columns are pixel numbers (784 total).
- Each value is the darkness of the pixel (1 to 255)

#### **Classes:**

0 - T-shirt/top, 1 – Trouser, 2 – Pullover, 3 – Dress, 4 – Coat, 5 – Sandal, 6 – Shirt, 7 -Sneaker, 8 – Bag, 9 - Ankle boot.

---

### **Model: Neural Network containing 2D-Convolution and Linear Transformation**

**2D-Convolution:** Applies a 2D convolution over an input signal composed of several input planes. In the simplest case, the output value of the layer with input size  $(N, C_{in}, H, W)$  and output  $(N, C_{out}, H_{out}, W_{out})$  can be precisely described as:

$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

where  $\star$  is the valid 2D cross-correlation operator,  $N$  is a batch size,  $C$  denotes a number of channels,  $H$  is a height of input planes in pixels, and  $W$  is width in pixels.

```
self.conv1 = nn.Conv2d(1, 6, 5)
self.conv2 = nn.Conv2d(6, 16, 5)
```

#### **Linear Transformation:**

Linear Transformation applies a linear transformation to the incoming data:  $y = x.A^T + b$

```
self.fc1 = nn.Linear(16*5*5, 120)
self.fc2 = nn.Linear(120, 84)
self.fc3 = nn.Linear(84, 10)
```

---

## Loss and Activation Functions

### Loss Function:

**Cross Entropy Loss** : Cross entropy quantifies the difference between two probability distribution. Our model predicts a model distribution of {p1, p2, p3 ....., p10} (where p1+p2+p3+.....+p10 = 1). We use cross-entropy to compare this with the true distribution {y1, y2, y3, ..., y10}

The losses are averaged across observations for each minibatch.

```
criterion = nn.CrossEntropyLoss()
```

$$\text{loss}(x, \text{class}) = \text{weight}[\text{class}] \left( -x[\text{class}] + \log \left( \sum_j \exp(x[j]) \right) \right)$$

### Activation Functions:

**ReLU** - Rectified linear unit. Applies the rectified linear unit function element-wise.

$$\text{ReLU}(x) = \max(0, x)$$

```
x = F.avg_pool2d(F.relu(self.conv1(x)), (2,2))
x = F.avg_pool2d(F.relu(self.conv2(x)), 2)
#
x = F.relu(self.fc1(x))
x = F.relu(self.fc2(x))
```

**Softmax** - Applies the Softmax function to an n-dimensional input Tensor rescaling them so that the elements of the n-dimensional output Tensor lie in the range (0,1) and sum to 1.

$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

```
return F.softmax(x,dim=1)
```

---

### Optimizer: Adam optimizer - *adaptive moment estimation*

Adam is an optimization algorithm that can used instead of the classical stochastic gradient descent procedure to update network weights iterative based in training data.

- Adam is a replacement optimization algorithm for stochastic gradient descent for training deep learning models.
- Adam combines the best properties of the AdaGrad and RMSProp algorithms to provide an optimization algorithm that can handle sparse gradients on noisy problems.
- Adam is relatively easy to configure where the default configuration parameters do well on most problems.

```
optimizer = optim.Adam(net.parameters(), lr=0.001)
```

## RESULTS

Experimental Prediction Comparisons: EPOCHS 25, LR 0.001, Saved in file **Output.txt**

Batch Size	Overall Accuracy	Class-wise Accuracy
32	88.72 %	T-shirt/top : 83 % Trouser : 96 % Pullover : 84 % Dress : 90 % Coat : 82 % Sandal : 96 % Shirt : 60 % Sneaker : 96 % Bag : 98 % Ankle boot : 93 %
64	88.94 %	T-shirt/top : 86 % Trouser : 94 % Pullover : 88 % Dress : 90 % Coat : 81 % Sandal : 98 % Shirt : 50 % Sneaker : 98 % Bag : 94 % Ankle boot : 91 %
<b>100</b>	<b>89.76 %</b>	<b>T-shirt/top : 91 %</b> <b>Trouser : 100 %</b> <b>Pullover : 85 %</b> <b>Dress : 92 %</b> <b>Coat : 87 %</b> <b>Sandal : 91 %</b> <b>Shirt : 75 %</b> <b>Sneaker : 96 %</b> <b>Bag : 94 %</b> <b>Ankle boot : 97 %</b>
256	87.76 %	T-shirt/top : 80 % Trouser : 100 % Pullover : 76 % Dress : 85 % Coat : 88 % Sandal : 100 % Shirt : 57 % Sneaker : 95 % Bag : 93 % Ankle boot : 100 %

### Best Execution Result: Batch Size = 100, LR = 0.001, Epochs 25

```
Device cuda 0
Properties: _CudaDeviceProperties(name='GeForce 840M', major=5, minor=0,
total_memory=4096MB, multi_processor_count=3)
Optimizer: Adam
Training Start.....
Batch size: 100
Epoch: 1/25, Approximated Loss: 1.129848426663125
Epoch: 2/25, Approximated Loss: 1.0628678306316999
Epoch: 3/25, Approximated Loss: 1.0472933776605091
Epoch: 4/25, Approximated Loss: 1.0382859022411237
Epoch: 5/25, Approximated Loss: 1.0331926876351063
Epoch: 6/25, Approximated Loss: 1.0284631376968758
Epoch: 7/25, Approximated Loss: 1.0228337952777695
Epoch: 8/25, Approximated Loss: 1.0192152504732668
Epoch: 9/25, Approximated Loss: 1.016727340132348
Epoch: 10/25, Approximated Loss: 1.0137242532972337
Epoch: 11/25, Approximated Loss: 1.012320290380252
Epoch: 12/25, Approximated Loss: 1.0093821302422084
Epoch: 13/25, Approximated Loss: 1.0061071088245355
Epoch: 14/25, Approximated Loss: 1.0055807673689396
Epoch: 15/25, Approximated Loss: 1.0032560717207135
Epoch: 16/25, Approximated Loss: 1.0003648341147342
Epoch: 17/25, Approximated Loss: 0.9998924694040924
Epoch: 18/25, Approximated Loss: 0.9982057388716853
Epoch: 19/25, Approximated Loss: 0.9966053588667699
Epoch: 20/25, Approximated Loss: 0.9962048587829606
Epoch: 21/25, Approximated Loss: 0.9952492450573656
Epoch: 22/25, Approximated Loss: 0.9936388998174107
Epoch: 23/25, Approximated Loss: 0.9926942509291903
Epoch: 24/25, Approximated Loss: 0.9914221831039787
Epoch: 25/25, Approximated Loss: 0.9909889172591675
Finished Training

```

---

```
The sample comparisons between ground truth vs the predicted labels

Ground Truth: T-shirt/top | Dress | Shirt | Sandal | Pullover | Pullover | Dress |
Trouser | T-shirt/top | Sandal
Predicted: T-shirt/top | Dress | Shirt | Sandal | Pullover | Pullover | Dress |
Trouser | T-shirt/top | Sandal

Accuracy of network: 89.76 %


```

---

```
Accuracies per class
Accuracy of T-shirt/top : 91 %
Accuracy of Trouser : 100 %
Accuracy of Pullover : 85 %
Accuracy of Dress : 92 %
Accuracy of Coat : 87 %
Accuracy of Sandal : 91 %
Accuracy of Shirt : 75 %
Accuracy of Sneaker : 96 %
Accuracy of Bag : 94 %
Accuracy of Ankle boot : 97 %
-----Plotting the predicted results of some random images-----
Figures saved in current working folder.
Done..

Process finished with exit code 0
```

Best Training Loss Values: Saved in file **Fmnist\_training\_loss\_a1.png**



Best Sample Predictions: Saved in files **0.png** and **1.png**



Predicted: Sandal

Correct: Sandal




Predicted: Sneaker

Correct: Sneaker



Predicted: Coat

Correct: Shirt



Predicted: Pullover

Correct: Pullover



Predicted: Sneaker

Correct: Sneaker



Predicted: Dress

Correct: Dress



Predicted: Coat

Correct: Coat



Predicted: Trouser

Correct: Trouser



Predicted: Bag

Correct: Bag



Predicted: Trouser

Correct: Trouser